

Unveiling the Power and Limitations of Unsupervised Node Embeddings

Charalampos E. Tsourakakis



Joint works with:



Dan Chanpuriya
UMass → UIUC →
Meta



Cameron Musco
UMass



Konstantinos
Sotiropoulos
BU → CMU
→ Meta

ML on Graphs is on the rise!

- Node Classification

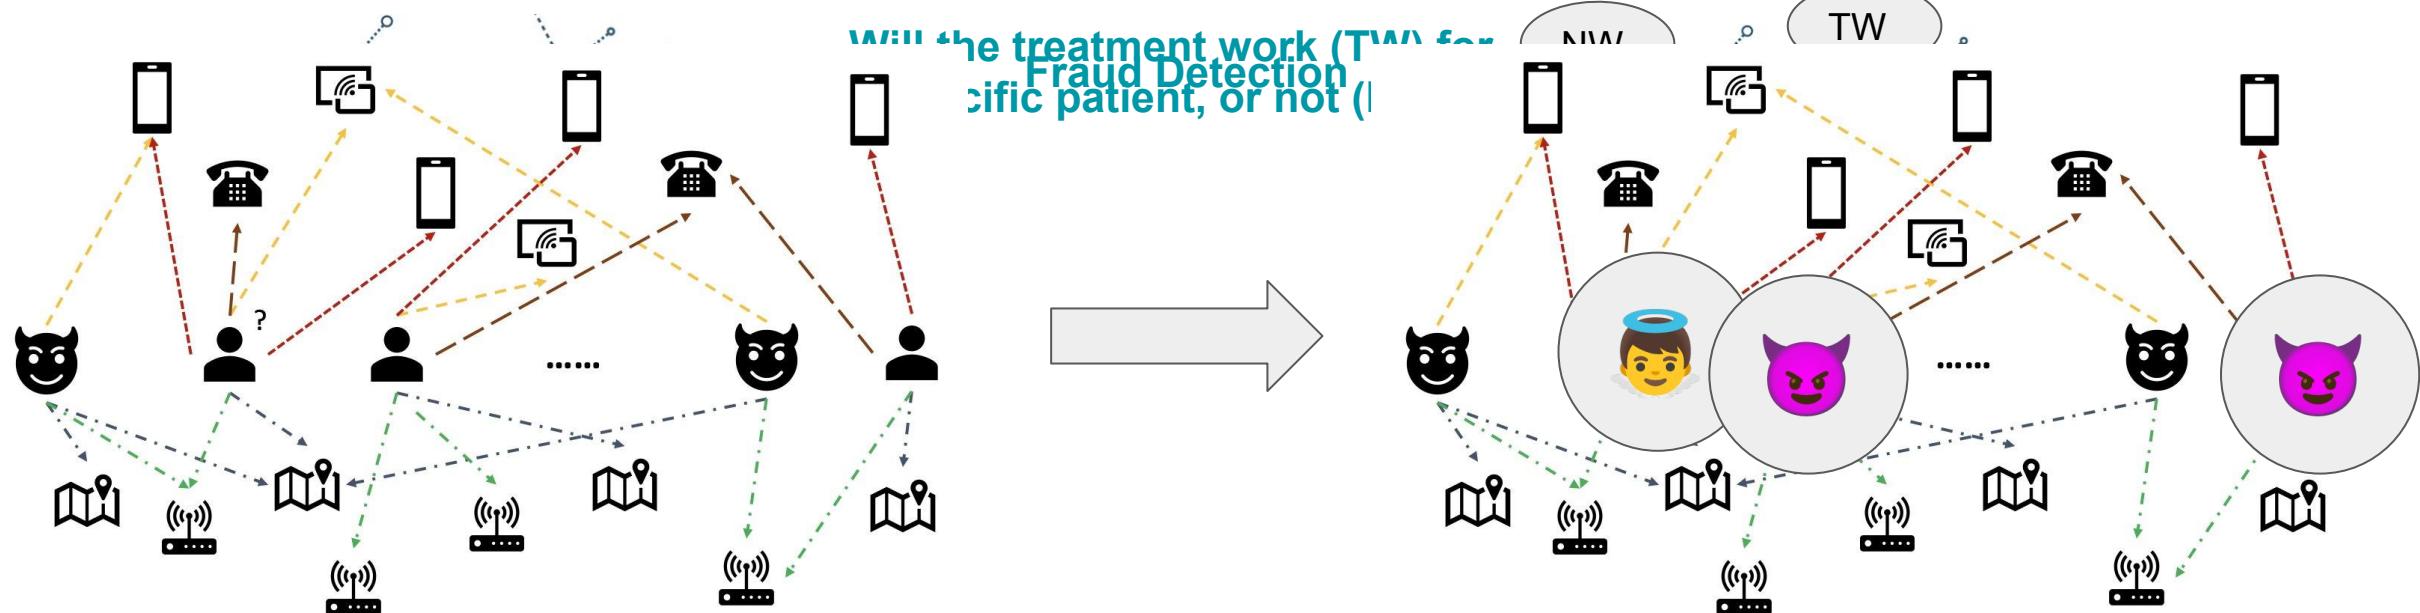
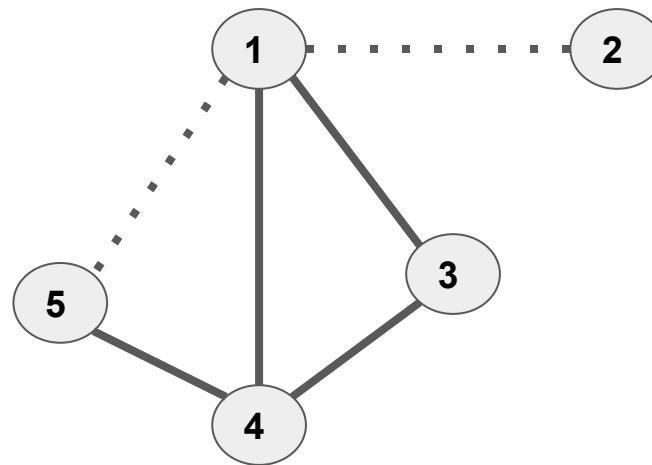


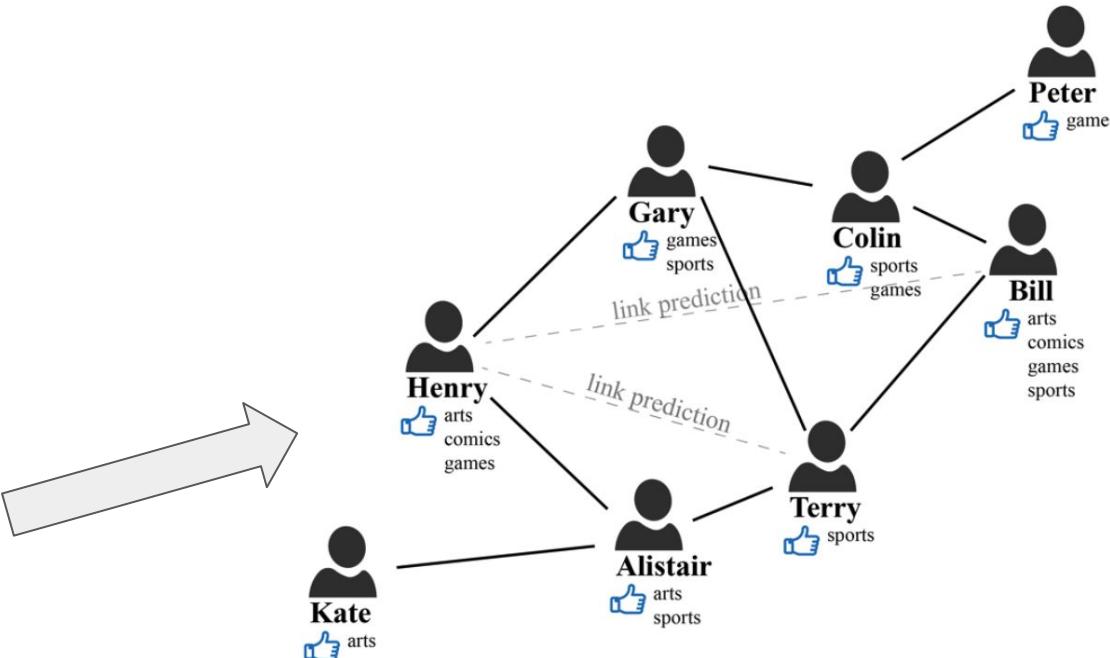
Image Source <https://davidsalak.com/>

ML on Graphs is on the rise!

- Link Prediction



- Friend Recommendation
- Entity Resolution
- Improve Marketing Strategies
- ...



Source: AWS Amazon

ML on Graphs is on the rise!

- Graph Generation
 - Model evolution of networks
 - Test null models (e.g. statistical significance of motifs)
 - Evaluate graph algorithms
 - Privacy
 - Discover new chemical and molecular structures



Erdős & Rényi, 1960



Watts-Strogatz, 1998



Chung-Lu, 2002

GraphVAE

Gilbert, 1959



Barabási-Albert-Bollobás-Riordan, 1999



GraphRNN

NetGAN

VGAE

What these tasks have in common?

- **(Unsupervised) Node Embeddings** are widely used as tools in all these tasks.
- Representation of nodes in a low-dimensional space $d \ll n$

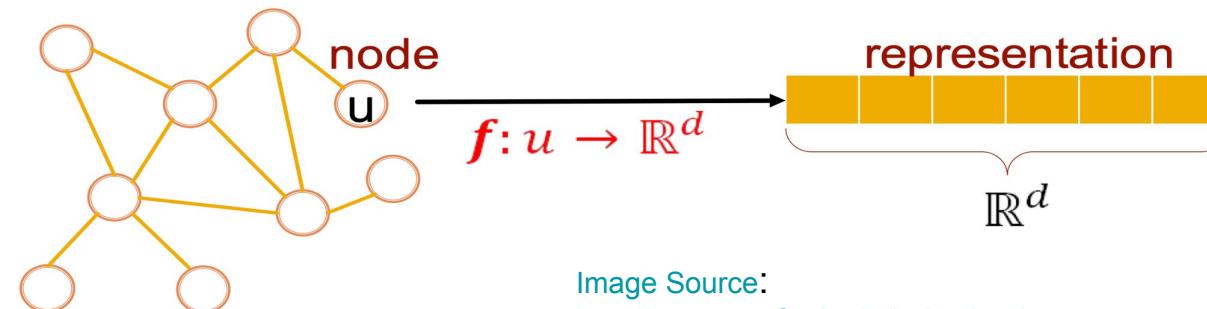
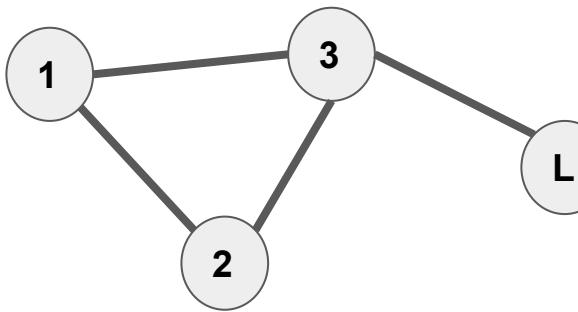


Image Source:
<http://snap.stanford.edu/index.html>

- **Why Unsupervised?**
 - Labels are not always available when we need them, e.g., in fraud detection. (Starnini et al., 2022)
 - Need to *better understand* unsupervised embeddings before we move to supervised.

How can we embed nodes?



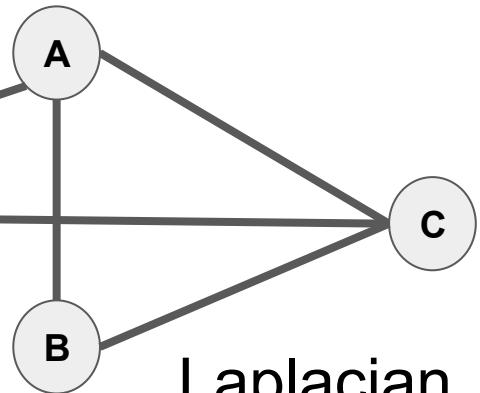
Adjacency Vectors

1	0	1	1	0	0	0	0	0
L	0	0	1	0	1	0	0	0
R	0	0	0	1	0	1	0	1
C	0	0	0	0	1	1	1	0

Structural
Features

(Degree, Triangles, Betweenness
Centrality)

2	1	0
2	0	0.57
3	1	0.57
3	2	0.12



Laplacian
Eigenvectors

0.09	0.28	0.45
0.38	-0.74	0.07
-0.65	-0.32	-0.23
-0.26	0.39	-0.40

Deep Learning

Word2Vec (Mikolov et al., 2003)

- **Skip-Gram model** (Mikolov et al., 2003) for word embeddings.
- **Input:** Collection of documents.

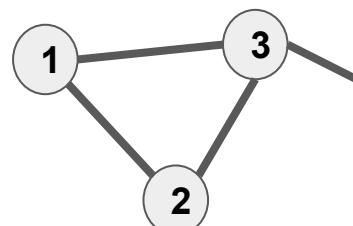


Image Source: Metropolitan Museum of Art

- **Target:** Embed words in a low-dimensional space & preserve semantic similarity.
- Words that co-appear within a window of size w in a corpus of sentences should have close representations.
- **Output:**
<https://projector.tensorflow.org/>

DeepWalk (Perozzi, Al-Rfou & Skiena, 2014)

- Document:



- Sentences: Random walks.

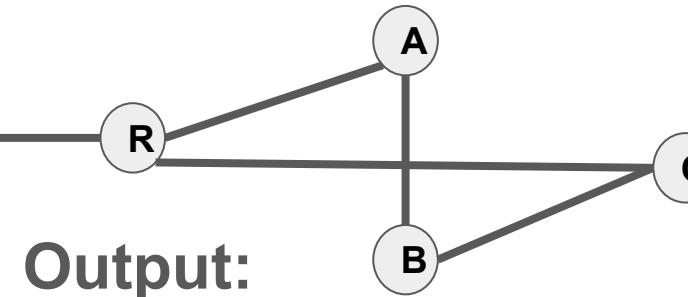
1, 3, 2, 3, L, R

3, 1, 2, 1, 2, 3

R, A, R, B, C, R

L, R, L, 3, 2, 1

...



Node2Vec (n2v) (Leskovec & Grover, 2016)

- Uses higher-order random walks.
- Explores graph in a BFS or DFS fashion (tunable parameters).

NetMF (Qiu, Dong, Ma, Li, Wang & Tang, 2018)

- Provides a unifying framework for many embedding methods as matrix factorization.
 - Among them DeepWalk and Node2Vec.
- DeepWalk implicitly factorizes the following matrix:

$$PMI(G, w) = M_w = \log \left(\frac{\text{vol}_G}{w} \sum_{r=1}^w (D^{-1} A)^r D^{-1} \right)$$

w: window size
D: diagonal matrix of degrees
 vol_G : Volume of graph

- Why PMI? Pointwise Mutual Information from Information Theory.

$$\log \left(\frac{p(x,y)}{p(x)p(y)} \right) \rightarrow \log \left(\frac{\text{vol}_G}{w} \sum_{r=1}^w (D^{-1} A)^r D^{-1} \right)$$

- Very important tool to analyze node embeddings through the lens of low-rank matrix factorization!

Node embedding methods

- Spectral Clustering (Alon & Milman, 1985, Shi & Malik, 2000, Ng et al., 2002, ...)
- IsoMap (Tenenbaum et al., 2000)
- Graphlets (Pržulj 2009), (Twittermancer)
- Locally Linear Embeddings (LLE) (Roweis & Saul, 2000)
- Laplacian Eigenmaps (Belkin & Niyogi, 2003)
- DeepWalk (Perozzi et al., 2014)
- node2vec (Grover & Leskovec, 2016)
- NetMF (Qiu et al., 2018)
- LINE (Tang et al., 2015)
- HARP (Chen et al., 2018)
- GraREP (Cao et al., 2015)
-

Deepwalk: Online learning of social representations

[B Perozzi](#), [R Al-Rfou](#), [S Skiena](#) - Proceedings of the 20th ACM SIGKDD ..., 2014 - dl.acm.org
... (DeepWalk) that learns **social representations** of a ... learn **social representations** with the following characteristics: • Adaptability - Real **social** networks are constantly evolving; new **social** ...
☆ Save ⚡ Cite Cited by 8857 Related articles All 22 versions

node2vec: Scalable feature learning for networks

[A Grover](#), [J Leskovec](#) - Proceedings of the 22nd ACM SIGKDD ..., 2016 - dl.acm.org
... **node2vec**, an algorithmic framework for learning continuous feature representations for nodes in networks. In **node2vec**, ... We demonstrate the efficacy of **node2vec** over existing state-of...
☆ Save ⚡ Cite Cited by 9323 Related articles All 24 versions

The impossibility of low-rank representations for triangle-rich complex networks

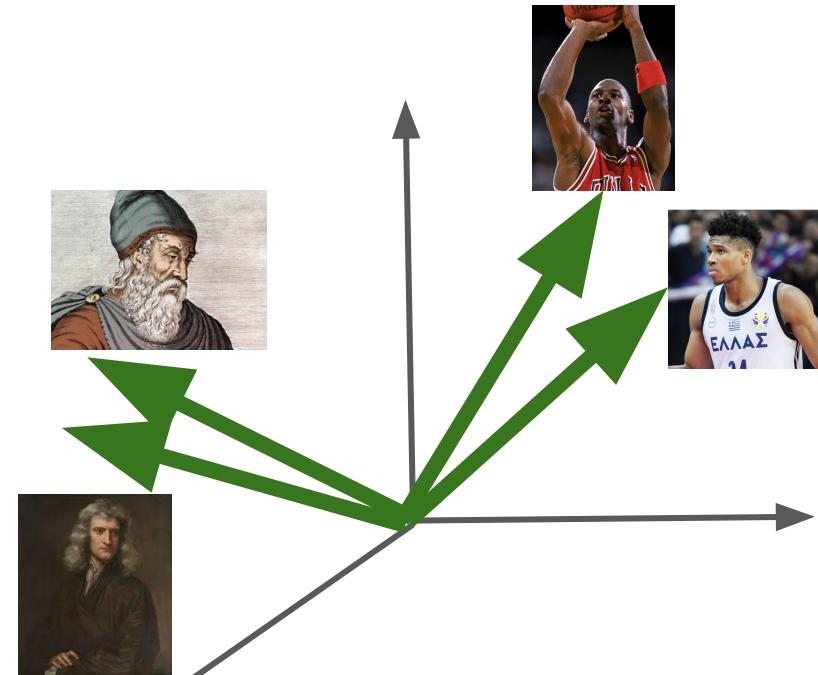
C. Seshadhri^{a,1}, Aneesh Sharma^b, Andrew Stolman^a, and Ashish Goel^c



Seshadhri et al. took a different research path.

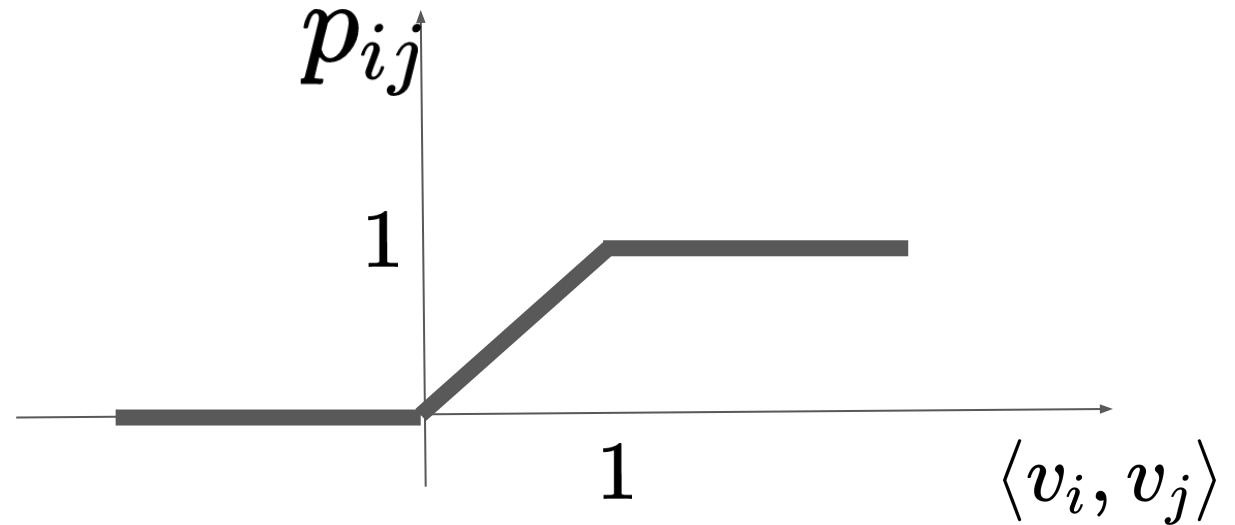
- What are the limitations of node embeddings?

Setting

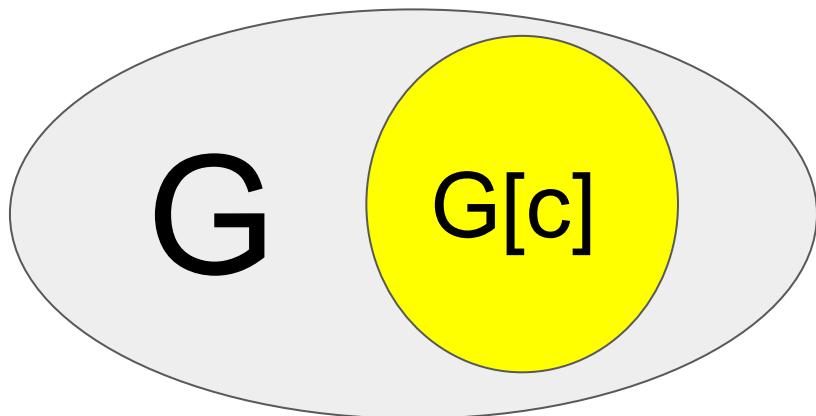


Setting: - Node i embedded as a vector v_i in d dimensions. In practice $d \ll n$.

- Connection probability of i, j by the truncated dot product (TDP)
$$p_{ij} = \max(0, \min(\langle v_i, v_j \rangle, 1))$$



(c, Δ) triangle foundation [Seshadhri et al.]



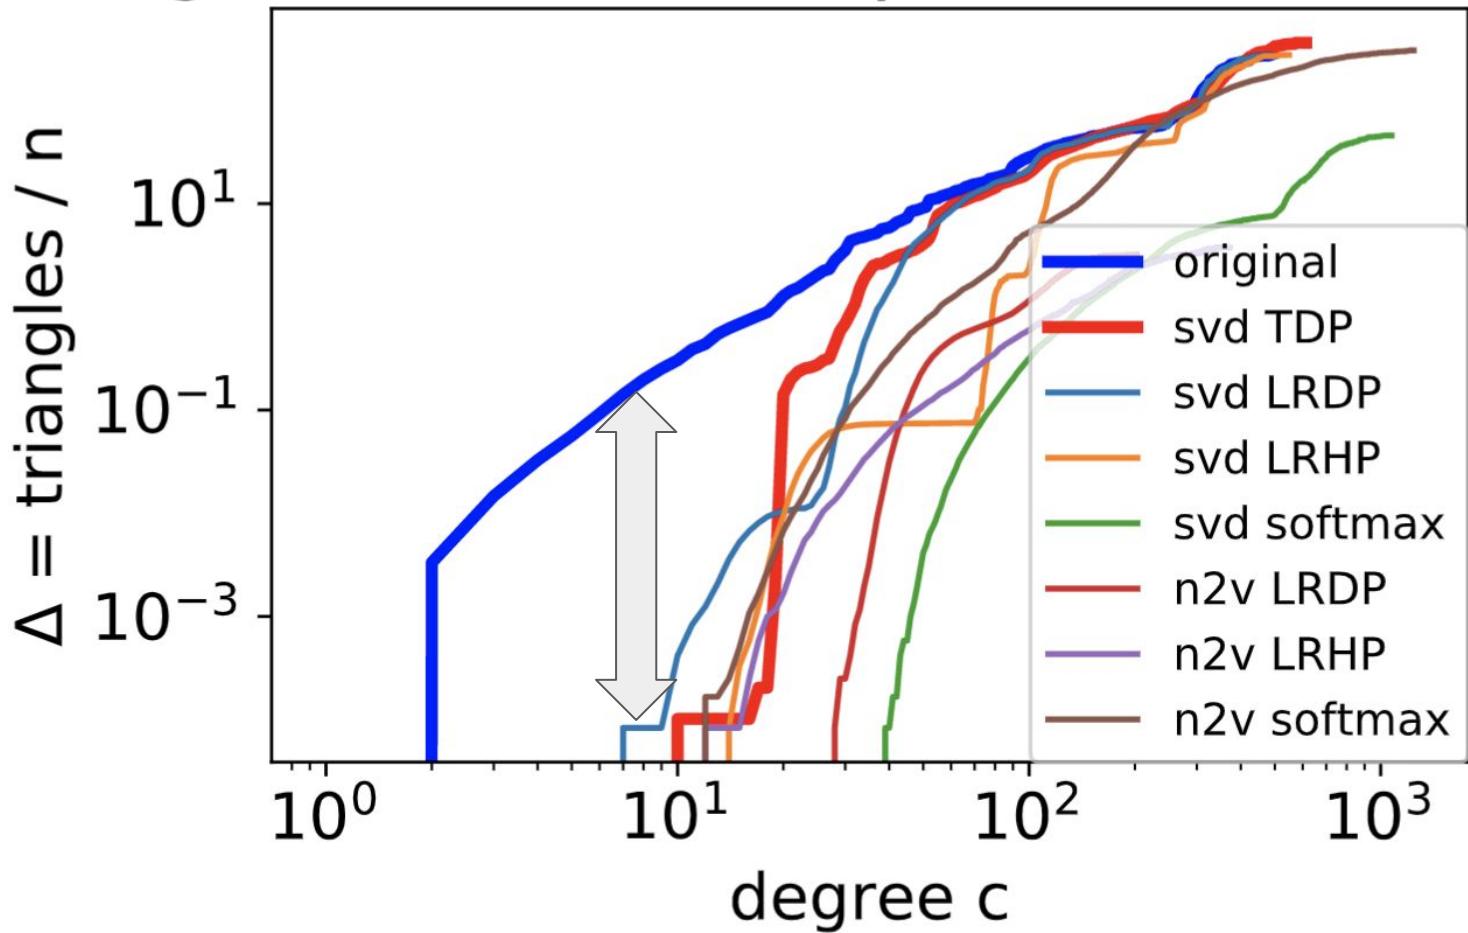
$G[c]$ = subgraph induced by all nodes whose degree is at most c

n = #number of nodes in G

$$\text{Define } \Delta(c) = \frac{\#\Delta s \text{ in } G[c]}{n}$$

Empirical observation: linear number of triangles among constant degree vertices.

Degree vs Δ of ca-HepPh and embeddings



DP: dot product

HP: Hadamard
product

T: truncated

LR: logistic
regression

svd: singular value
decomposition
n2v: node2vec

Theorem

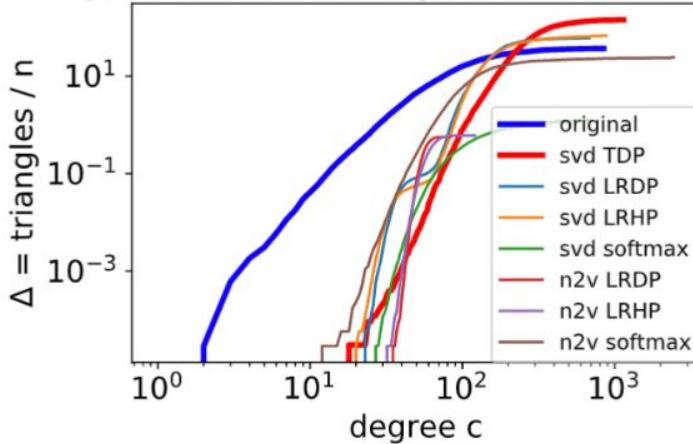
Fix $c > 4$, $\Delta > 0$. Suppose the expected number of triangles in $G[c]$ is at least Δn . Then, the dimensionality d is at least

$$\min\left(1, \text{poly}\left(\frac{\Delta}{c}\right)\right) \frac{n}{\log^2 n}$$

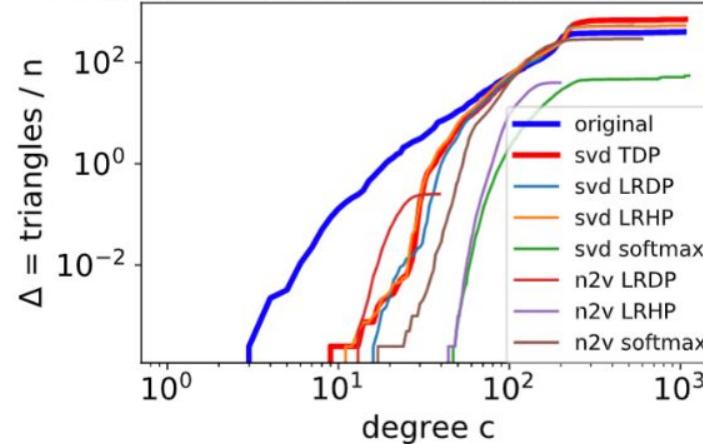
- In other words, graphs generated from low-dimensional embeddings cannot contain many triangles only on low-degree vertices

Theory agrees with practice

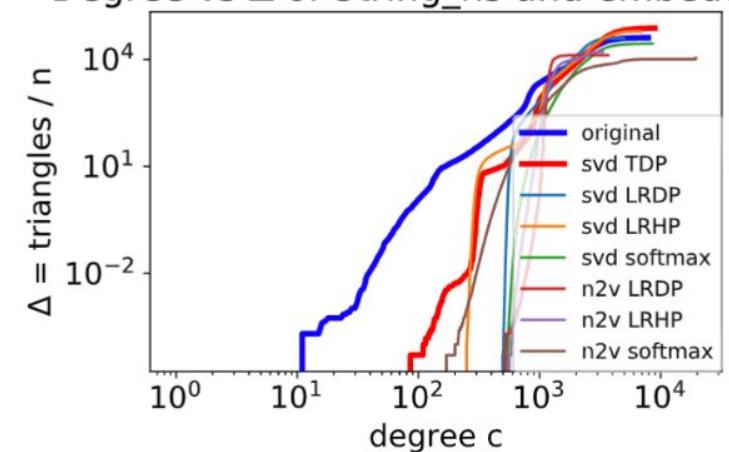
Degree vs Δ of cit-HepPh and embeddings



Degree vs Δ of Facebook and embeddings



Degree vs Δ of String_hs and embeddings



- 100 dimensions and 100 graph samples, plotting the best among them. Still, the gap is evident!
- Equivalently, inherent limitations in representing triangle dense graphs if the embedding is given by $X \in \mathbb{R}^{n \times d}$ and the expected adjacency matrix is given

$A = \text{clip}(XX^T)$, where clip clips entries to [0,1].

Possibility Result

- (Seshadhri, Sharma, Stolman & Goel, PNAS 2020):

“Low-dimensional embeddings can not generate graphs that have many triangles among low-degree nodes (sparse and triangle rich)”

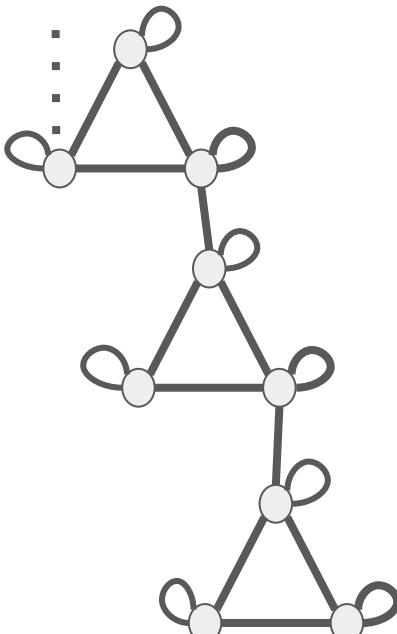
- (Chanpuriya, Musco, Sotiropoulos & T, NeurIPS 2020):

“Low-dimensional embeddings can not only generate graphs that have many triangles among low-degree nodes (sparse and triangle rich), but exactly encode the graph!”

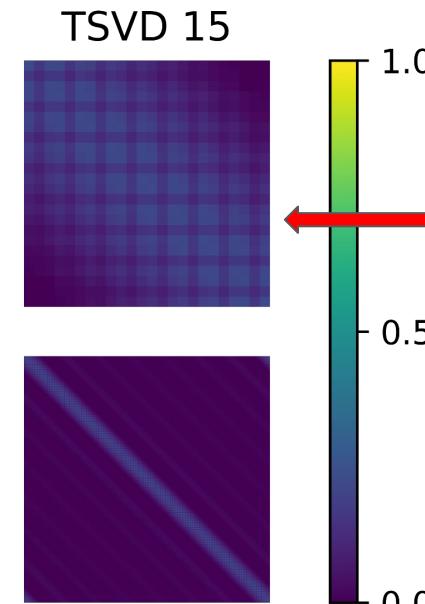
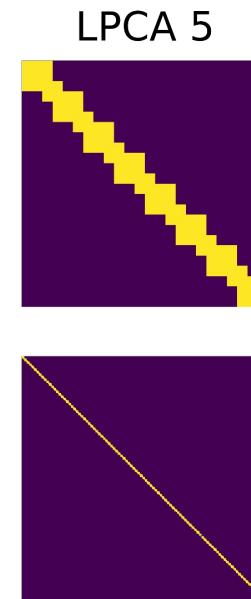
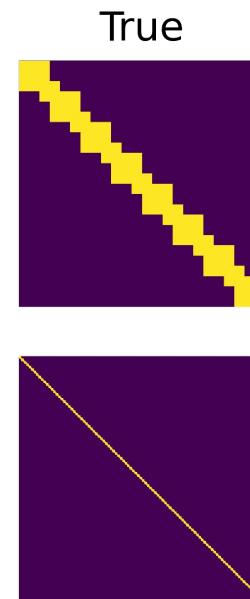
- How can both statements be true?

Reconstructing a sparse and triangle-dense graph

Graph

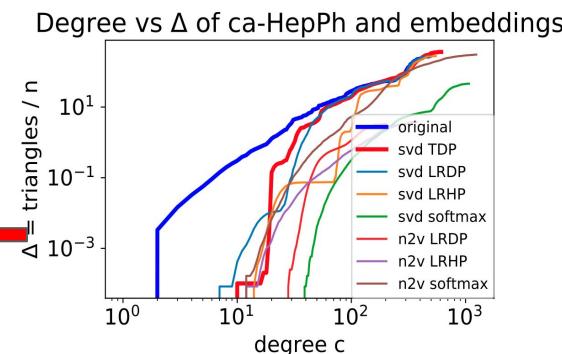


Reconstructions from Embeddings



LPCA
Logistic PCA

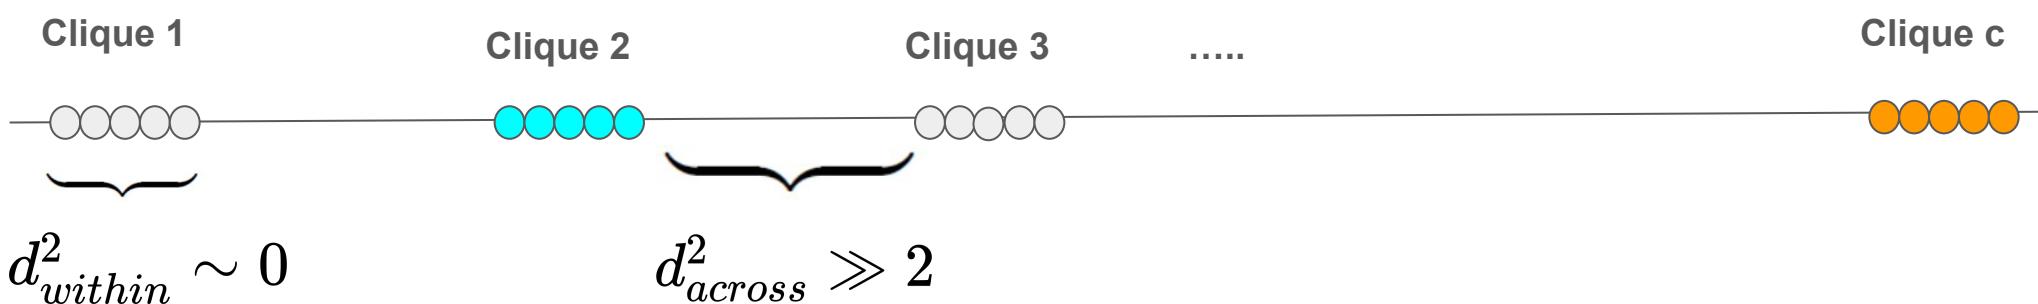
TSVD
Truncated SVD



Relaxing the PSD constraint (2 embeddings per node)

- **Observation:** By relaxing the PSD constraint of the dot-product model we can exactly factorize a sparse and triangle dense graph using only 3 dimensions.

Example: Let A be the adjacency matrix of a union of n/c cliques.



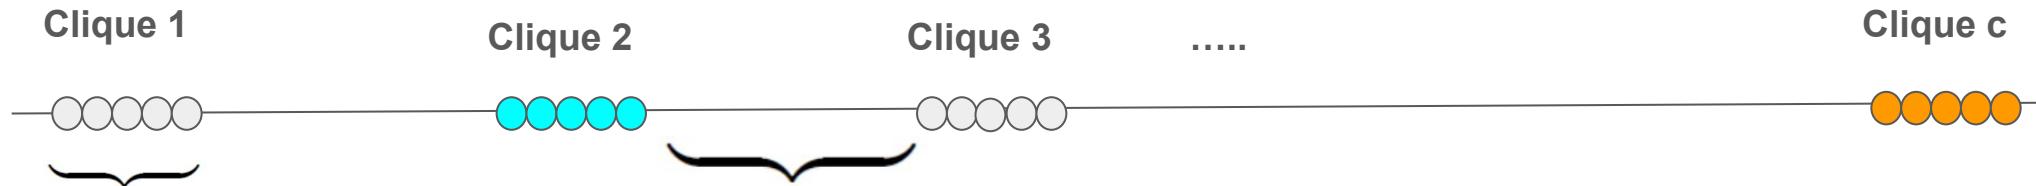
The matrix D_{dist} of squared euclidean distances has rank $\leq d+2$ ($d=1$ in our case).

Consider $\tilde{A} = 2 - D_{dist}$.

\tilde{A} has rank 3 and $A = clip(\tilde{A}) = clip(XY^T)$ where $X, Y \in \mathbb{R}^{n \times 3}$

Relaxing the PSD constraint (2 embeddings per node)

In the 1D case it is easy to see the rank-3 factorization of



$$d_{within}^2 \sim 0$$

$$d_{across}^2 \gg 2$$

$$D_{dist} = \begin{bmatrix} 0 & & & \\ & 0 & & \\ & & \ddots & (x_i - x_j)^2 \\ & & (x_i - x_j)^2 & 0 \\ & & & & 0 \end{bmatrix}$$

Define

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}, y = \begin{bmatrix} x_1^2 \\ x_2^2 \\ \vdots \\ x_n^2 \end{bmatrix}, \mathbb{1} = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$

$$\mathbb{1}y^T + y\mathbb{1}^\top - 2xx^T = \begin{bmatrix} x_1^2 & \cdots & x_n^2 \\ \vdots & \ddots & \vdots \\ x_1^2 & \cdots & x_n^2 \end{bmatrix} + \begin{bmatrix} x_1^2 & \cdots & x_1^2 \\ \vdots & \ddots & \vdots \\ x_n^2 & \cdots & x_n^2 \end{bmatrix} - 2 \begin{bmatrix} x_1^2 & \cdots & x_1x_n \\ \vdots & \ddots & \vdots \\ x_1x_n & \cdots & x_n^2 \end{bmatrix} = D_{dist}$$

Connection with sign-rank

$$A = \text{clip}(XX^T) \rightarrow A = \text{clip}(XY^T)$$

1. Matrix is PSD. Relax this by considering XY^T .
2. We apply a function that clips entries to $[0, 1]$. Think of A as $\{-1, 1\}$ instead of $\{0, 1\}$.
We are interested in a low-rank factorization that ***matches A sign-wise***.

Definition [Sign-rank]: The sign rank of a real matrix A , with entries $\{-1, 1\}$ is the least rank of a matrix B , such that $A_{ij} \cdot B_{ij} > 0$ for all i, j .

Studied in TCS due to connections with circuit complexity (Razborov & Sherstov, 2010), (Bun & Thaler, 2016), communication complexity (Alon et al., 1985), (Linial et al., 2007), learning theory (Alon et al., 2016).

Key Theorem

[CMST, NeurIPS 2020]

For any graph G with maximum degree c , there exist factors $X, Y \in \mathbb{R}^{n \times (2c+1)}$ of dimension $2c+1$, such that $A = \text{clip}(XY^T)$

Proof sketch: Let $V^{nx(2\Delta+1)}$ be the Vandermonde matrix evaluated at $1,..,n$

$$V^{n \times (2c+1)} = \begin{bmatrix} 1^0 & 1^1 & \dots & 1^{2c} \\ 2^0 & 2^1 & \dots & 2^{2c} \\ 3^0 & 3^1 & \dots & 3^{2c} \\ 4^0 & 4^1 & \dots & 4^{2c} \\ \vdots & \vdots & \ddots & \vdots \end{bmatrix}$$

Notice that for any vector x^{2c+1} , $(Vx)[t] = \sum_{i=0}^{2c} x_i t^i$

Key Theorem

- Let a_i be the i -th row of the adjacency matrix A .
 - It has at most c ones.
- For each node i , we find a polynomial of degree $2c$ which is positive at all integers t with $a_i(t)=1$ and negative at all t with $a_i(t)=0$.
 - Straight-forward interpolation
 - The coefficients X give us $A = \text{sign}(VX^T)$

Corollary: preferential attachment graphs have maximum degree $O(\sqrt{n})$ with high probability, so it admits sublinear embedding.

Follow-up result: $O(\alpha^2)$ dims suffice, $\alpha=\text{arboricity}$



Logistic PCA (LPCA)

- We also provide an algorithm that learns embeddings $X, Y \in \mathbb{R}^{n \times k}$ such that $A = \text{clip}(XY^T)$ using gradient-descent on the following logistic loss:
$$\text{loss} = - \sum_{i,j} \sigma(\tilde{A}_{ij} \cdot [XY^T]_{ij}),$$
 where σ is the sigmoid.
- We can exactly factorize real-world networks using very few dimensions.

Dataset	# Nodes	# Edges	$2\Delta + 1$	Exact Empirical Factorization Dimension
Pubmed	19 581	43 847	343	48
ca-HepPh	11 204	117 634	983	32
BlogCatalog	10 312	333 983	7 985	128
Citeseer	3 327	4 552	199	16
Cora	2 708	5278	337	16

Are the embeddings of DeepWalk exact?

- LPCA can exactly encode the input graph.
- It is not clear what DeepWalk (or NetMF) embeddings encode.

Question:

“What information of the input graph
do some of the most popular
modern embeddings encode?”

Our contribution

- Invert Node Embeddings:
 - We invert NetMF embeddings by learning graphs that have similar embeddings to the input graph.
 - Study the properties of these graphs to understand what is encoded in low-dimensions.
 - How is this information relevant to downstream ML tasks?
- **Problem Formulation**

Given an embedding algorithm $\mathcal{E} : \mathcal{G} \rightarrow \mathbb{R}^{n \times d}$ and the embedding $\mathcal{E}(G)$ for some $G \in \mathcal{G}$, produce $\tilde{G} \in \tilde{\mathcal{G}}$ such that $\|\mathcal{E}(\tilde{G}) - \mathcal{E}(G)\|$ is small. How close are G, \tilde{G} in terms of, e.g., edges, degrees, triangles, community structure?

Algorithm 1 for $d=n$, $w \rightarrow \infty$ - Solve a linear system!

- In the limit, as $w \rightarrow \infty$, using a result from Chanpuriya & Musco (2020) we know that:

$$M_{\infty}^{n \times n} = \text{vol}_G D^{-1/2} \left(L_{normalized}^{\dagger} - I \right) D^{-1/2} + \mathbf{1}\mathbf{1}^T$$

- Solving a linear system yields the pseudo-inverse
- Cons, in practice:
 1. w is never set to be infinite, instead typically a small value
 2. The embeddings in practice are low dimensional, not full/high dimensional.

Algorithm 2 - Practical PMI optimization algorithm

- For typical settings, $w=10$ and low-rank, we employ an optimization-based approach.
- Given the low-rank embedding $M_{w,d}$ of a graph G , our algorithm finds a graph \tilde{G} of same expected volume with a similar embedding.

$$\min_{\tilde{G}} \left\| PMI(\tilde{G}, w) - M_{w,d} \right\|_F^2$$

- We apply this algorithm to invert embeddings of popular benchmark graphs in node classification tasks.

Experimental Setting

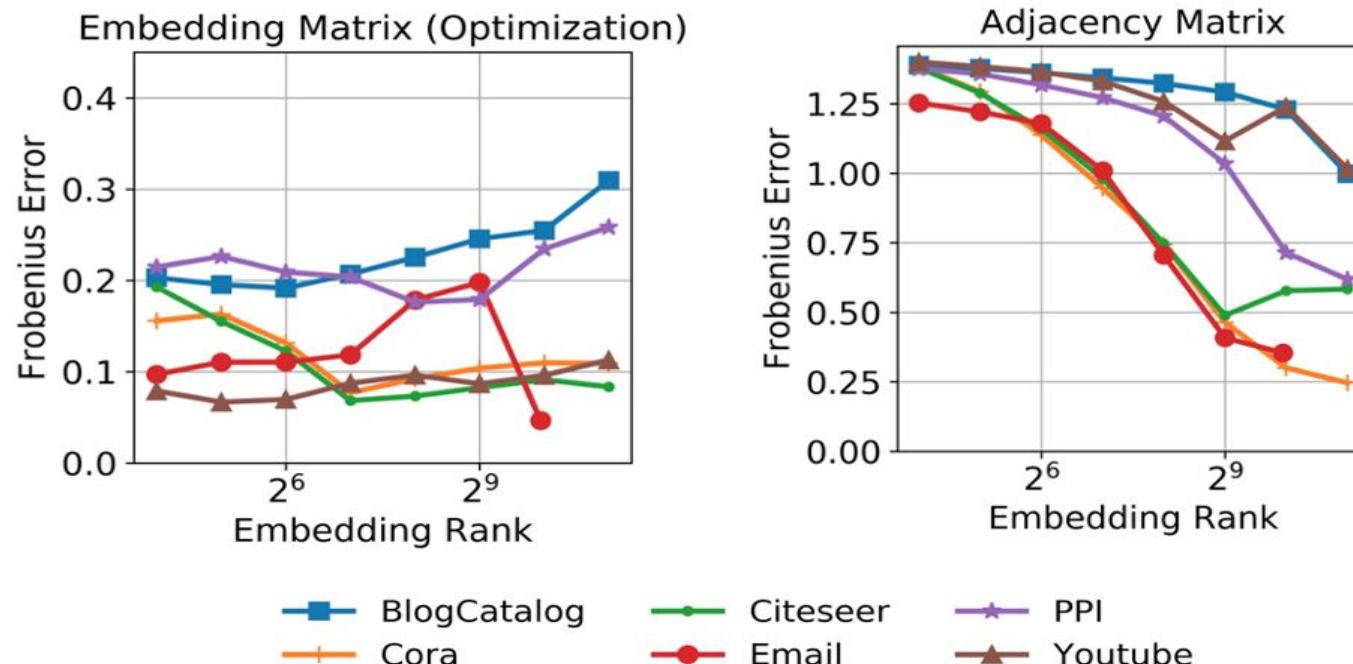
- We use 6 popular benchmark datasets in node classification (#labels): BlogCatalog (39), Cora (7), Citeseer (6), Email (42), PPI (50), YouTube (20)
- We use PyTorch, auto-differentiation and SciPy for the optimization algorithm.
- Reconstruction errors for embeddings and adjacency matrices \tilde{M}_r and \tilde{A}_r (rank r) are reported using the (relative) Frobenius Error w.r.t. the input graph M and A.

$$\frac{\|X - \tilde{X}\|_F}{\|X\|_F}$$

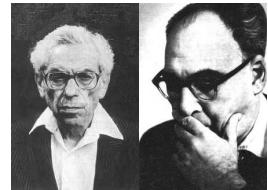
- We evaluate embeddings in node classification tasks. We use one-vs-rest logistic regression classifier and report the Micro-F1 score (global F1 score).

Results

- In the y-axis we report the Frobenius Error as we vary the embedding dimensions (x-axis).
 - 1) We observe the embedding error is small (and not sensitive to dimension).
 - 2) Despite having similar embeddings, reconstructed networks are very different in low ranks.



Graph Generative Models



Erdős & Rényi, 1960



Watts-Strogatz, 1998



Chung-Lu, 2002

Gilbert, 1959

Barabási-Albert-Bollobás-Riordan, 1999



GraphVAE

GraphRNN

NetGAN

VGAE

Our contributions [CMST NeurIPS '21]:

- Embeddings used to form an edge independent model are limited, there is an inherent trade-off between expressivity and memorization.
- *“Study graph generative models under the lens of how capable they are to mimic the original network while generating “edge diverse” samples.”*

Graph Generative Models

More recently we see deep-learning based models.

- Given an input graph G , ...
- ... **learn** a distribution and sample new graphs *similar* to G .
 - NetGAN (Bojchevski et al., 2018)
 - VGAE (Kipf & Welling, 2016)
 - GVAE (Simonovski & Komodakis 2018)
 - GraphRNN (You et al., 2018)
 -

Evaluation of these models

- We want generated graphs to be “similar” to the input graph.
- Match statistics of the input graph:
 - Degree distribution
 - Triangle count
 - Path length
 - Clustering coefficient (transitivity)
 - ...
- We can simply return the input graph, but...
 - We want a graph similar, but not identical to the input graph.
 - We also want samples to be edge-diverse.
- **Target:** Match graph statistics, while generating edge-diverse samples.

Motivation

- Complex Deep-Learning based models, like NetGAN (Bojchevski et al., 2018) or CELL (Rendsburg et al, 2020) learn a matrix $P \in [0, 1]^{n \times n}$ and sample edges independently.
- They stop training before sampled graphs overlap a lot with the input graph.
- Generated graphs usually have few triangles!

GRAPH	TRIANGLE COUNT
CORA-ML	2,802
NETGAN	(54% EO)
CELL	(53% EO)

GRAPH	TRIANGLE COUNT
WEB-EDU	4491
NETGAN	(53% EO)
CELL	(54% EO)

- Is this a coincidence? Namely, are there some limitations of edge independent models?

A few Preliminaries

- Our focus: Undirected, unweighted & labeled graphs ($V=[n]$).
- Let \mathcal{U}, \mathcal{V} be distributions over $\binom{n}{2}$ -dimensional vectors.

Inner Product: $\langle \mathcal{U}, \mathcal{V} \rangle := \mathbb{E}_{\mathbf{u} \sim \mathcal{U}, \mathbf{v} \sim \mathcal{V}} [\mathbf{u} \cdot \mathbf{v}]$

- Suppose \mathcal{A} is a distribution over adjacency matrices of undirected graphs on n nodes.
- Let \mathcal{F} be the distribution that always returns the complete graph K_n .

We define the expected volume as $\text{Vol}(\mathcal{A}) := \langle \mathcal{A}, \mathcal{F} \rangle$

Overlap

- Suppose \mathcal{A} is a distribution over adjacency matrices of labeled, undirected graphs on n nodes. We define the overlap of \mathcal{A} as:

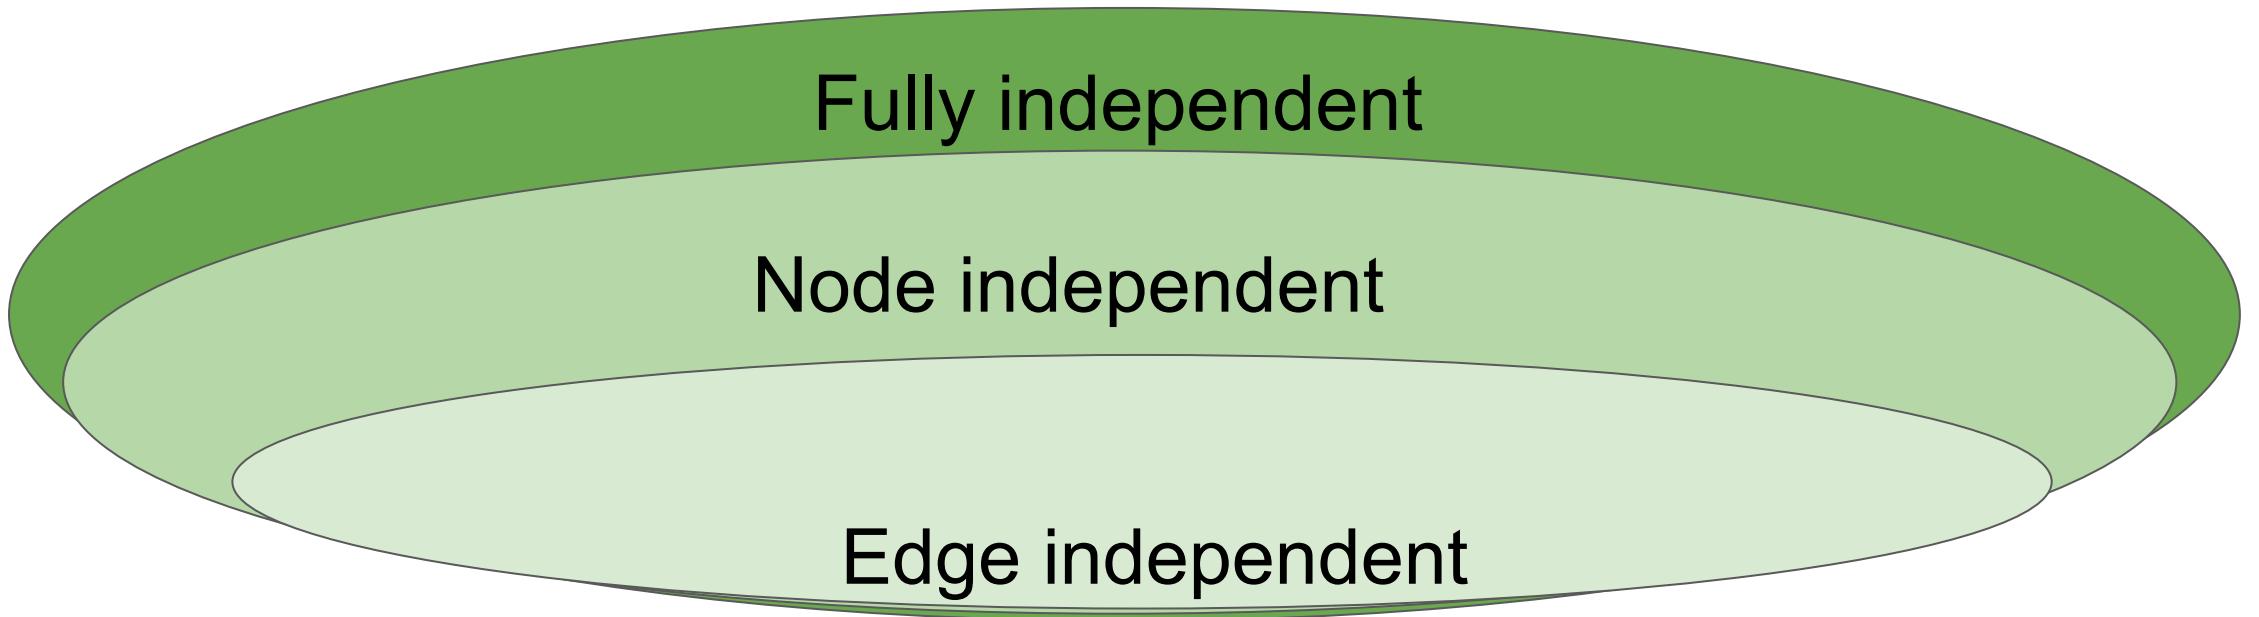
$$\text{Ov}(\mathcal{A}) := \frac{\langle \mathcal{A}, \mathcal{A} \rangle}{\text{Vol}(\mathcal{A})}$$

- **Meaning:** Expected fraction of edges that will overlap among two samples from the same distribution.
- Connection between Overlap and Volume (from Cauchy-Schwartz):

Lemma 3: $\text{Vol}(\mathcal{A}) \leq \binom{n}{2} \cdot \text{Ov}(\mathcal{A})$

“Question: For a given overlap, how well can different graph generative models mimic properties of the input graph, e.g., number of triangles?”

A Hierarchy of Graph Generative Models



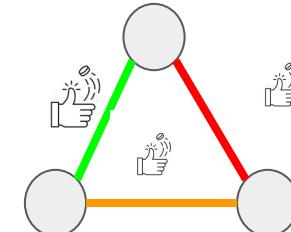
The node independent model instead of embedding each node to a vector, it gives to each node an independent distribution over embedding vectors.

Examples

Coin tosses are independent!

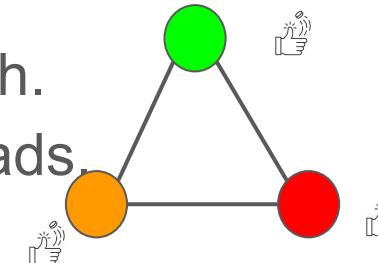
- **Edge Independent Model $\mathcal{A}_{E.I.}$.**

- Toss a coin with probability p for each edge in the graph.
- Probability of an edge = $Ov(\mathcal{A}_{E.I.}) = p$.
- Number of triangles $\approx n^3 \cdot Ov(\mathcal{A}_{E.I.})^3$



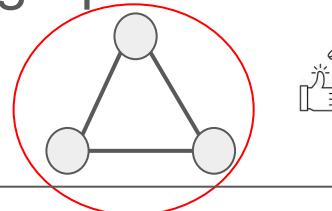
- **Node Independent Model $\mathcal{A}_{N.I.}$.**

- Toss a coin with probability \sqrt{p} for each node in the graph.
- Create a clique among all nodes that coin tosses are heads.
- Probability of an edge = $Ov(\mathcal{A}_{N.I.}) = p$.
- Number of triangles $\approx n^3 \cdot Ov(\mathcal{A}_{E.I.})^{3/2}$



- **Fully Dependent Model $\mathcal{A}_{F.D.}$.**

- With probability p create a complete graph, or empty graph otherwise.
- Probability of an edge = $Ov(\mathcal{A}_{F.D.}) = p$.
- Number of triangles $\approx n^3 \cdot Ov(\mathcal{A}_{F.D.})$



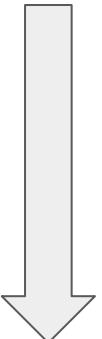
A Hierarchy of Graph Generative Models

- **Edge Independent Models** [Edges are independent]
A distribution \mathcal{A} over symmetric adjacency matrices \mathbf{A} , such that for some $\mathbf{P} \in [0, 1]^{n \times n}$, $\mathbf{A}_{ij} = \mathbf{A}_{ji} = 1$ with probability P_{ij} independently for all i, j .
- **Node Independent Models** [Models dependencies at node level]
A distribution \mathcal{A} over symmetric adjacency matrices $A = [XY^T > 0] \in \{0, 1\}^{n \times n}$ where \mathbf{X}, \mathbf{Y} are random matrices and the rows of $[XY]$ are mutually independent.
($[M > 0]$ sets all positive entries to 1, and 0 otherwise.)
- **Fully Dependent Models** [Models any kind of dependencies]
Any distribution \mathcal{A} over symmetric adjacency matrices.

Summary of Results

- Turns out these examples provide (asymptotically) tight bounds for the expected number of triangles, graphs generated from these models can have.

Model	Upper Bound in #triangles
Edge Independent Model	$n^3 \cdot \text{Ov}(\mathcal{A})^3$
Node Independent Model	$n^3 \cdot \text{Ov}(\mathcal{A})^{3/2}$
Fully Dependent Model	$n^3 \cdot \text{Ov}(\mathcal{A})$



Recall $\text{Ov}(\mathcal{A}) \leq 1$.

Putting these ideas in practice.

- We design a simple baseline for each model in the hierarchy.
- Desiderata:
 - Easy to range overlap.
 - Be able to reproduce (in the limit) the input graph.
- **CCOP (E.I):** Edge independent model.
Convex combination of input graph and random graph that preserves degree sequence on expectation.



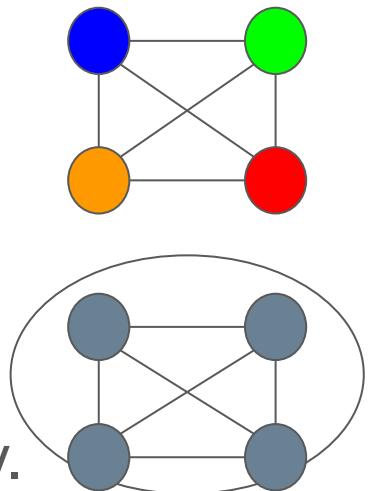
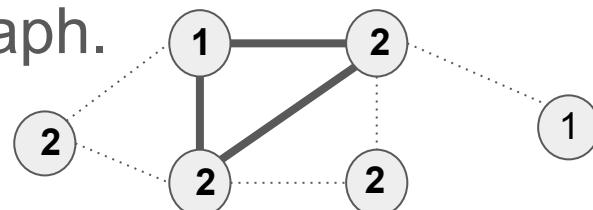
C-L $\alpha = 0$



True Graph $\alpha = 1$

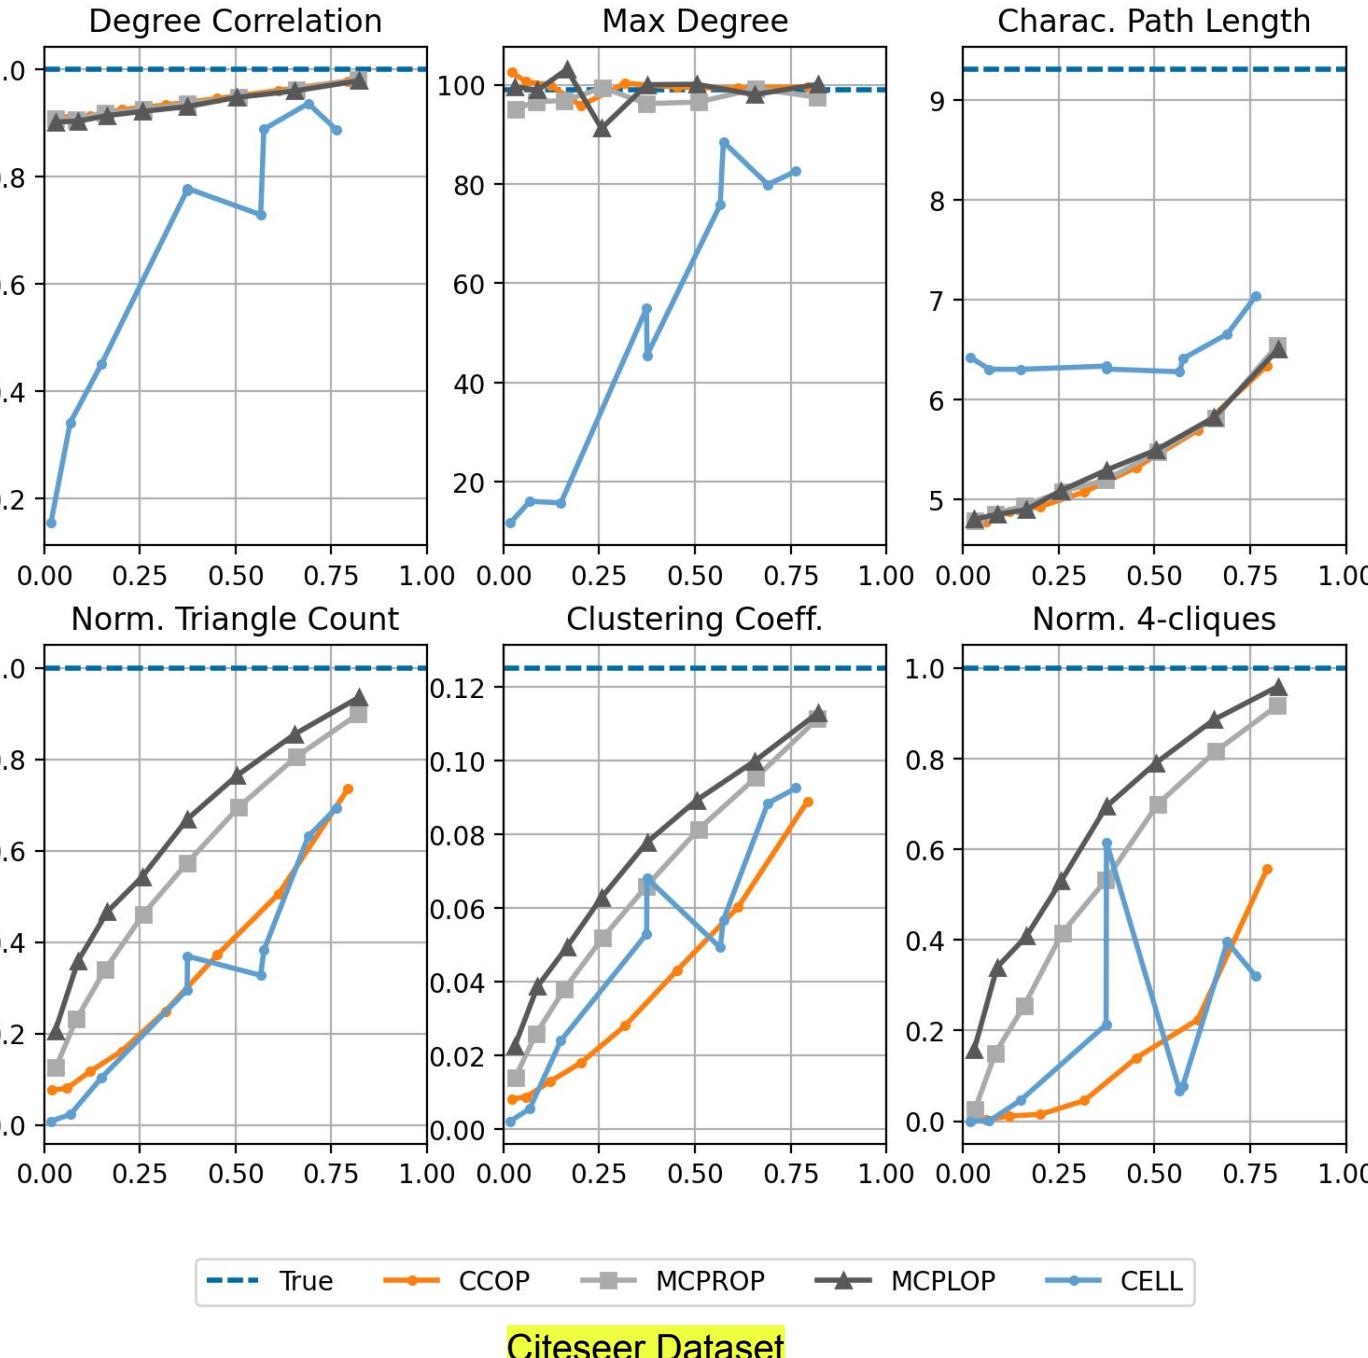
Putting these ideas in practice.

- Common idea for models with dependency.
- List all maximal cliques of the input graph.
 - **MCPROP (N.I.)**
For each maximal clique, retain each node with probability p , independently.
 - **MCPLOP (F.D.)**
Include each maximal clique with probability p , independently.
- Fit a random graph model that preserves degree sequence in expectation on the residual graph.



Experimental Evaluation

- We use 5 publicly available datasets:
 - Citeseer, Cora, PolBlogs, PPI, Web-Edu
- Datasets range from more sparse (e.g., Citeseer with 2.1K nodes and 7.3K edges) to more dense (e.g., PolBlogs with 1.2K nodes and 33.4K edges).
- We compare against CELL, a variant of the popular NetGAN method that applies then GAN framework on random walks.
 - Following the authors of CELL, we perform early-stopping to prevent from simply reproducing the input graph.
- We observe our baselines are competitive to the CELL method, and usually outperform it in matching counts of dense subgraphs in low overlap.



x-axis: Overlap,
y-axis: graph statistics

- Our baselines, especially those introducing dependencies, achieve higher triangle density in low overlap compared to CELL.
- By design, they also do well in matching degree sequence.
- CELL does better in average path length.

Importance of some of our key results

- Exact encoding of a graph with low dimensional embeddings
 - Sign rank as a key tool for understanding exact compression.
 - Motivates the use of mixtures of embeddings.
- We can go back from embeddings to graphs
 - Studying salient and non-salient characteristics.
 - Raises question about the privacy of node embeddings.
- We provide a unique lens of studying graph generative models.
 - Concept of overlap for refined random graph model evaluation.
 - We give a theoretical explanation to empirical findings on edge independent models.
 - Novel and simple baselines that are competitive to state-of-the-art models.

Open Questions

- How can we convert our deeper understanding to efficient practical methods?
 - Mixtures of embeddings?
- What is a good choice for the number of dimensions given the dense subgraph structure of the graph?
- How do we extend our findings to unlabeled graphs and directed graphs?

References

- (Broder et al., 2000) Broder, Andrei, et al. "Graph structure in the web." *Computer networks* 33.1-6 (2000): 309-320.
- (Mikolov et al., 2013) Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." *NIPS 2013*.
- (Perozzi et al., 2014) Perozzi, Bryan, Rami Al-Rfou, and Steven Skiena. "Deepwalk: Online learning of social representations." *KDD 2014*.
- (Grover & Leskovec, 2016) Grover, Aditya, and Jure Leskovec. "node2vec: Scalable feature learning for networks." *KDD 2016*.
- (Qiu et al., 2018) Qiu, Jiezhong, et al. "Network embedding as matrix factorization: Unifying deepwalk, line, pte, and node2vec." *WSDM 2018*.
- (Starnini et al., 2021) Starnini, Michele, et al. "Smurf-based anti-money laundering in time-evolving transaction networks." *ECML PKDD 2021*.
- (Erdős & Rényi, 1960) Erdős, Paul, and Alfréd Rényi. "On the evolution of random graphs." *Publ. Math. Inst. Hung. Acad. Sci* 5.1 (1960): 17-60.
- (Gilbert, 1959) Gilbert, Edgar N. "Random graphs." *The Annals of Mathematical Statistics* 30.4 (1959): 1141-1144.

References (2)

- (Watts & Strogatz, 1998) Watts, Duncan J., and Steven H. Strogatz. "Collective dynamics of 'small-world' networks." *nature* 393.6684 (1998): 440-442.
- (Barabási and Albert, 1999) Barabási, Albert-László, and Réka Albert. "Emergence of scaling in random networks." *science* 286.5439 (1999): 509-512.
- (Chung & Lu, 2002) Chung, Fan, and Linyuan Lu. "Connected components in random graphs with given expected degree sequences." *Annals of combinatorics* 6.2 (2002): 125-145.
- (Simonovsky & Komodakis, 2018) Simonovsky, Martin, and Nikos Komodakis. "Graphvae: Towards generation of small graphs using variational autoencoders." *ICANN*, 2018.
- (Kipf & Welling, 2016) Kipf, Thomas N., and Max Welling. "Variational graph auto-encoders." *arXiv preprint arXiv:1611.07308* (2016).
- (You et al., 2018) You, Jiaxuan, et al. "Graphrnn: Generating realistic graphs with deep auto-regressive models." *International conference on machine learning*. PMLR, 2018.
- (Bojchevski et al., 2018) Bojchevski, Aleksandar, et al. "Netgan: Generating graphs via random walks." *International conference on machine learning*. PMLR, 2018.

References (3)

(Rendsburg et al, 2020) Rendsburg, Luca, Holger Heidrich, and Ulrike Von Luxburg. "Netgan without gan: From random walks to low-rank approximations." *International Conference on Machine Learning*. PMLR, 2020.